

Task Coordinates for Few-Shot Diffusion Adaptation: Results Against the Experimental Plan

Anonymous authors
Paper under double-blind review

17 September 2026

Typeset with the fallback preamble: iclr2026-conference.sty was not found.

Abstract

We report the results of the current experimental stage for Meta-Diffusion, a method that adapts a frozen denoising diffusion probabilistic model to a new data domain by fitting a low-dimensional task coordinate rather than by updating network weights. The report is organised block by block against the governing experimental plan, so that every item the plan specifies appears here with either its result or the reason it was not run. The plan fixes the criterion for the main claim in advance: the claim is supported only if knowledge of the correct source task, combined with a learned source-to-target relation, outperforms *both* transfer that knows only the relation *and* few-shot modelling that uses only target-domain data. On the distributional metric the plan names as headline, that criterion is met in the low-target-data regime and only there. With a single target image the method improves on relation-only transfer by $+0.025 \pm 0.008$ and on target-only modelling by $+0.018 \pm 0.009$; the advantage over target-only modelling has vanished by twenty target images. Of the ten deliverables the plan fixes, six are complete and four are deferred. Two earlier readings are withdrawn, both traceable to a measurement instrument blind to the quantity under test.

1 SCOPE, CRITERION AND NOTATION

This report follows the plan’s own structure and numbering. Sections 2 through 6 correspond one to one with the plan’s Stage 0 through Stage D; Section 7 reports against its cross-cutting requirements and Section 8 against its list of final deliverables. Each block states what the plan asked, then either what was run and what it returned, or why it was not run.

The criterion. The plan states the criterion for the main claim without qualification: the claim is supported only if the correct source task *plus* the learned source-to-target relation outperforms *both* relation-only transfer *and* target-only few-shot modelling, in the low-target-data regime. Any other outcome is to be reported as what it is.

The model. The backbone is a denoising diffusion probabilistic model: a network trained to predict the noise added to an image at a randomly chosen point along a fixed corruption schedule, so that running the schedule backwards generates samples. Generation uses the denoising diffusion implicit model sampler with the stochasticity parameter set to zero, which makes sampling deterministic given the initial noise.

The coordinate and the transport map. A set encoder r_ψ maps a set of images to a task coordinate $z \in \mathbb{R}^k$. Writing z_S for the coordinate encoded from source-domain data and c for a descriptor of the source-to-target relation, the transport map produces the predicted target coordinate $\tilde{z}_T = z_S + \Delta_\gamma(z_S, c)$. The map takes no target image as input; target images enter only through an optional refinement step that minimises the denoising loss over the coordinate alone, with the network frozen.

Notation, as the plan defines it. M_S is the number of source-domain images the encoder sees when producing z_S . K_T is the number of target-domain images available at adaptation time. Feature-wise linear modulation, abbreviated below as the modulation arm, is the standard alternative to an additive basis: it scales and shifts intermediate feature maps as a function of the coordinate instead of adding a learned direction to the score.

Two instruments. The first is a transformation statistic: three fixed convolutions measuring mean absolute Laplacian response, per-image standard deviation and residual above a local average. It reads which image transformation was applied and nothing else. The second is the sliced Wasserstein distance in pixel space, which compares two sets of images by projecting both onto many random directions, sorting each projection and comparing quantiles. The plan names the second as the headline distributional metric. Section 3.1 shows that the choice between them is not cosmetic.

Statistics. Every comparison is a paired difference computed episode by episode and reported with a 95% confidence interval; an interval excluding zero is marked with an asterisk. An interval containing zero means the sign is not determined at this sample size. It does not establish equality, which would require a separately declared margin of indifference that none of these comparisons carries.

2 STAGE 0 — PRECONDITIONS

2.1 0.1 Integrate the three branches

What the plan asked. Merge the measurement branch and the modulation branch into the baseline, so that `scripts/gen_strategies.py`, `models/film_unet.py` and `baselines/film_conditioning.py` exist in one tree and the test suite passes.

Not run, and why. The merge was not performed. In the tree that produced every number in this report, `scripts/gen_strategies.py` is present and both modulation-arm files are absent. This is the single unmet precondition with downstream consequences: it blocks A2, and through A2 it blocks deliverable 3. Nothing else in the plan depends on it, which is why the remainder of Stage A could proceed.

2.2 0.2 Verify the split file

What the plan asked. Confirm the task-family split file against a fixed checksum before the first run, and pass it explicitly on every command.

Result. Verified. The file hashes to `2ae9b81b...47708`, matching the value the protocol fixes. Every command reported here passed the path explicitly.

2.3 0.3 Recover or retrain the modulation checkpoint

What the plan asked. Obtain or retrain `film128_step50000.pt`, since checkpoints do not travel with the branch.

Not run, and why. The checkpoint is absent from this machine and was neither recovered nor retrained. Together with 0.1 this is the second reason A2 could not run; the two are independent, so both would have to be resolved.

2.4 0.4 Add the missing evaluation primitives

What the plan asked. Five pieces of code: a source-budget flag on the generation script; nested source subsampling so that a source-budget sweep varies set size alone; a relation-only transport condition; a within-relation shuffled source; and the kernel Inception distance.

Result. Four of the five were written and are exercised by the results below. The source budget flag drives A4; nested subsampling makes that sweep paired rather than confounded; the relation-only condition and the within-relation shuffle are two of the three controls in A1. **The kernel Inception distance was not implemented.** Its absence is the reason Section 9 can make no statement about absolute generative quality, and it should be read as a gap in this report rather than as a decision.

2.5 0.5 Create the output directory

What the plan asked. Write results under the hand-off tree, never under the artefacts tree, whose ignore rules would discard them silently.

Table 1: Training and adaptation configuration, read from the checkpoint’s own stored configuration and from the adaptation budget object shared by every strategy.

Training		Adaptation (meta-test)	
Backbone	U-Net, 128 base channels	Refinement steps J	25
Channel multipliers	(1, 2, 2), 2 residual blocks	Coordinate rate η_z	10^{-2}
Coordinate dimension k	16	Prior weight β_0	1.0, divided by K_T
Set encoder pooling	mean	Noise draws per step	32
Transport hidden width	64	Sampling	implicit, $\eta = 0$
Relation descriptor d_c	8		
Diffusion steps	1000, cosine schedule		
Optimiser	AdamW, 2×10^{-4}		
Warmup, gradient clip	500 steps, 1.0		
Moving-average decay	0.999		
Loss weights λ_T, λ_{tr}	1.0, 1.0		
Coordinate penalty λ_z	10^{-4}	Reverse steps	50
Source budget at training M_S	64	Samples per cell	96
		Episodes per cell	60

Result. Done, and observed throughout.

3 STAGE A — CIFAR-100, GENERATION-LEVEL EVIDENCE

All of Stage A uses one checkpoint of the three-transformation CIFAR-100 domain-shift task family — Gaussian blur, additive Gaussian noise and contrast reduction at severity 3 — trained for 50,000 steps and read with exponentially moving averaged weights. Table 1 records the configuration.

A departure from the written configuration, now corrected. The configuration file recorded 32 noise draws per refinement step only after this stage began. The field previously read 16 and was never passed to the object that consumes it, so every number produced before the correction was in fact generated with that object’s own default of 32. The field was changed to match what had been running rather than the reverse. Any earlier document quoting 16 is wrong.

A change that altered conclusions. Refinement originally drew its noise from an unseeded global stream, so two strategies compared within one episode refined against different noise: an unpaired difference inside an otherwise paired comparison. A per-episode stream, rewound for each strategy, was introduced part-way through the stage. Two readings moved when it was, and both are identified where they occur.

3.1 A1 — Source-dependence controls

What the plan asked. Four conditions on the three-transformation setup — the correct source coordinate, a within-relation shuffled source, the mean source coordinate and relation-only — with the distributional metric as headline and paired intervals throughout. The plan adds an instruction: if relation-only matches the full method, say so explicitly and narrow the source-informed claim, because that outcome is reportable rather than a failure.

What was run. 180 episodes, twenty classes by three transformations by $K_T \in \{1, 5, 20\}$, nine conditions, 96 samples per cell.

Result. The method beats all three source controls on the headline metric, in all nine cells, every interval clear of zero (Table 2). The outcome the plan warned about did not occur: relation-only is comfortably beaten.

The criterion is met, with a boundary. At a single target image the method beats relation-only transfer and target-only modelling simultaneously, which is what the plan requires. The advantage over target-only modelling survives to five target images and is gone by twenty. The claim is therefore explicitly bounded: source information is worth having

Table 2: A1, paired differences on the headline metric. Positive values favour the full method. The final row repeats the comparison against the mean source coordinate under the transformation statistic, where every interval contains zero.

Full method against	$K_T = 1$	$K_T = 5$	$K_T = 20$
Mean z_S	$+0.0301 \pm 0.0072^*$	$+0.0346 \pm 0.0085^*$	$+0.0338 \pm 0.0080^*$
Shuffled z_S	$+0.0416 \pm 0.0125^*$	$+0.0453 \pm 0.0149^*$	$+0.0436 \pm 0.0144^*$
Relation-only	$+0.0254 \pm 0.0078^*$	$+0.0285 \pm 0.0088^*$	$+0.0295 \pm 0.0087^*$
Target-only	$+0.018 \pm 0.009^*$	$+0.011 \pm 0.007^*$	-0.003 ± 0.007
Mean z_S , transformation statistic	$+0.017 \pm 0.033$	-0.008 ± 0.030	$+0.011 \pm 0.028$

when target evidence is scarce, and its value decays as target evidence accumulates. That boundary is consistent with the task family, in which a single target image already reveals which transformation was applied.

The two instruments disagree, and the disagreement is the finding. Under the transformation statistic the comparison against the mean source coordinate contains zero at every budget. That instrument reads which transformation was applied; the within-relation mean preserves the transformation exactly and destroys only task identity. The instrument is blind by construction to precisely the quantity the control removes, so a null reading from it carries no information about that quantity. An earlier reading of this panel concluded that the task-specific content of the coordinate went unused. That conclusion was drawn from the blind instrument and is withdrawn.

The gap to the oracle is real. The method remains behind the oracle by $+0.026$, $+0.027$ and $+0.017$ at the three budgets, every interval excluding zero. Earlier statements that the method was indistinguishable from the oracle came from the loss-shaped instrument and are likewise withdrawn.

3.2 A2 — Additive basis against the modulation arm, at generation

What the plan asked. Extend the existing loss-level comparison to generation, matched on backbone, latent dimension and budget.

Not run, and why. Blocked by two independent unmet preconditions: the branch integration of 0.1, which leaves the modulation implementation outside this tree, and the missing checkpoint of 0.3. Neither is a scientific obstacle and neither was resolved during the stage. The arm itself is built and has been measured on the denoising loss; only the generation-level comparison is outstanding. This is the sole incomplete item of Stage A.

3.3 A3 — Target-scarcity sweep

What the plan asked. Sweep K_T and name the crossing point at which target-only modelling catches up.

What was run. 300 episodes at $K_T \in \{1, 2, 5, 10, 20\}$.

Result. The method beats reusing the source coordinate at every budget, the advantage decaying monotonically from $+1.140 \pm 0.108$ to $+0.229 \pm 0.059$ with every interval clear of zero. This is the largest and most robust effect in the stage, an order of magnitude above every other comparison.

The crossing point, as the plan asks it to be named. There is no crossing. The correct statement is a convergence: the advantage over target-only modelling is confined to the smallest budgets and the two become indistinguishable thereafter. An earlier run had read the twenty-image budget as the method *losing*, -0.017 ± 0.044 , and that was written up as a crossing. The larger run reads $+0.008 \pm 0.036$: the sign reversed and both intervals contain zero, so the method was never shown to lose. That write-up is retracted.




Figure 1: A4. Six source budgets from 16 to 300 images, one target image, 60 episodes each, one checkpoint throughout. Left: absolute distance per strategy. Right: paired differences with 95% intervals. Nothing moves.

3.4 A4 — Source-abundance sweep

What the plan asked. Fix one target image and sweep the source budget over $\{16, 32, 64, 128, 256, 300\}$, using the nested subsampling of 0.4 so that set identity is not confounded with set size. The plan states its expected shape: improvement with increasing source evidence, then saturation.

What was run. All six budgets, 60 episodes each, one checkpoint throughout.

Result. Saturation from the first point measured, with no improvement phase at all (Figure 1). Across a nineteenfold range the comparison against reusing the source coordinate moves 0.0002, that against target-only modelling 0.0019 and that against the oracle 0.0011 — each smaller than the half-width of its own interval at either end. Sixteen source images already buy everything this transport map extracts from a source set. Reported as what it is, per the plan’s instruction.

The sweep’s internal check. Two strategies in the panel never consult the source coordinate, and they return identical values at all six budgets. That is the sweep’s own confirmation that it varies only what it claims to vary, and it passes.

Scope. Every point is the same checkpoint, trained once at $M_S = 64$, evaluated with a different number of source images. The sweep measures how much source evidence a trained transport map requires at deployment. It does not measure how that map would differ had it been trained under a different source budget; see Section 7.3.

3.5 A5 — Transport-structure ablation

What the plan asked. Four maps from the source coordinate to the predicted target coordinate under one generation protocol: direct reuse, a constant relation vector, a linear map and the residual network the method uses.

What was run. Four models trained separately under identical conditions, compared episode by episode over 180 episodes each.

Result. Table 3. The identity is beaten decisively. The linear map, with nineteen times fewer parameters, is indistinguishable from the residual multilayer perceptron at every budget.

Table 3: A5, paired against the residual multilayer perceptron (6824 parameters). Positive values favour it. The final row is the noise floor: the spread across the same four runs of a strategy that never consults the transport map, and therefore ordinary run-to-run variation.

Variant	Parameters	$K_T = 1$	$K_T = 5$	$K_T = 20$
Identity	0	$+0.0242 \pm 0.0080^*$	$+0.0256 \pm 0.0089^*$	$+0.0246 \pm 0.0087^*$
Constant relation vector	96	$+0.0082 \pm 0.0051^*$	$+0.0092 \pm 0.0052^*$	$+0.0109 \pm 0.0055^*$
Linear map	352	$+0.0022 \pm 0.0049$	$+0.0004 \pm 0.0048$	$+0.0036 \pm 0.0050$
Run-to-run floor	—	0.0076	0.0036	0.0135

The floor, and what it permits us to claim. These are four separately trained models, so a difference between them mixes the structure of the map with ordinary variation between training runs. The panel prints its own floor: a strategy that never consults the transport map spreads by 0.0076 to 0.0135 across the same four runs, against a transport spread of 0.0082 to 0.0109 once the identity is excluded. Read against that floor, the defensible conclusion is narrow. The map must be more than the identity, by a margin comfortably above the floor at every budget. The constant relation vector’s margin sits inside the floor. On these numbers the linear map is the structure to prefer, and the claim that the non-linear map earns its place is not supported.

4 STAGE B — GAUSSIAN-MIXTURE POSITIVE CONTROL

What the plan asked. A cheap positive control on two-dimensional Gaussian mixtures, where the ground truth is analytic, together with a sweep of the distance between source and target tasks.

What was run. Five task families spanning that distance, 32 tasks each.

Result. The mechanism behaves as designed. The paired improvement from using a task’s own coordinate rather than the coordinate shared by all tasks runs from $+0.1428$ to $+0.1571$ across the five families, every interval excluding zero, and the task-specific share of the gain is 1.0: on analytic mixtures essentially the whole benefit of the coordinate is specific to the task rather than a constant offset. The coordinate-loss profile is minimised at the task’s own coordinate in all five families, which is the correctness check the panel exists to perform.

Figure 2 places this beside the corresponding measurement on images. The two loss scales are not comparable, so the figure plots the depth of the well as a percentage of the loss at the shared mean coordinate, which is. The well is 26.1% deep on mixtures and 1.6% deep on CIFAR-100, and both are minimised in the right place with the wrong direction rising monotonically. The mechanism on images is therefore neither absent nor misdirected. It is shallow, and an earlier reading that described it as flat was overstating.

What the sweep did not produce. The relatedness sweep was built to vary the distance between source and target tasks, and it does not move the task-specific term: the five readings are $+0.1539$, $+0.1428$, $+0.1536$, $+0.1571$ and $+0.1533$, from the smallest perturbation through to unrelated task pairs. The dial turns; the quantity it was built to move does not. We record this as an open question rather than a footnote.

5 STAGE C — DEMOGRAPHIC SCARCITY EXPERIMENT

5.1 C0 — Feasibility check (hard gate)

What the plan asked. A same-day check, gating the rest of Stage C, on whether the dermatological dataset annotated by skin phototype can be assembled.

Result. A qualified pass. The annotation table is fully available: 16,577 rows over 114 conditions. The images are not. One of the two source archives holds 12,631 of those rows, 76.2% of the dataset, and its domain fails to resolve for public resolvers as well as locally, with the delegation itself reported as unreachable; a sample of thirty of its links returned no image bytes at all. The second archive is fully live, and thirty of thirty sampled links returned images.




Figure 2: The coordinate-loss profile on two-dimensional Gaussian mixtures (left) and on CIFAR-100 domain shift (right). The horizontal axis walks the coordinate from the shared mean at $s = 0$ through the task’s own coordinate at $s = 1$ and out to twice that displacement. Both are minimised at $s = 1$; the depth of the well differs by roughly sixteenfold.

Grouping prototypes I–III against IV–VI, chosen from the observed counts rather than in advance, the reachable archive yields 16 conditions above the plan’s floor. Only 8 of those can fund a twenty-image support set together with a held-out query set, which gives a five/two/one split of tasks. **A single held-out test task cannot carry a paired interval**, and paired intervals are what every other result in this report rests on. The reachable subset is therefore scaffolding for C1–C3, not a basis for the headline result of C4. Restoring the remaining three quarters of the images depends on a request from the dataset’s authors control.

5.2 C1 — Domain module and splits

What the plan asked. A domain module and a split builder that splits the semantic tasks themselves into train, validation and test, so that the relation rather than the task becomes the object of generalisation.

Not run, and why. Not started. It is buildable today on the reachable subset and is not blocked; it was not reached within the stage.

5.3 C2 — Frozen evaluation predictors

What the plan asked. A frozen semantic consistency predictor and a phenotype consistency predictor, each reporting its own held-out ceiling beside every verdict.

Not run, and why. Not started, and dependent on C1.

5.4 C3 — Optional sanity check

What the plan asked. Verify on a few tasks that the model produces a measurable target-population shift when target evidence is abundant. Explicitly a sanity check, not a gate.

Not run, and why. Not started, and dependent on C1 and C2.

5.5 C4 — The scarcity sweep

What the plan asked. Six arms across $K_T \in \{1, 2, 5, 10, 20\}$, reported in the form of target-data equivalence — for example, the method at two target images matching target-only modelling at ten — which the plan prefers to quoting a small metric improvement.

Not run, and why. Blocked by the qualification on C0, not merely unreached. The reachable data support eight usable conditions, which cannot fund a task split with more than one held-out test task, and the equivalence statement the plan asks for cannot be made without a paired interval on held-out tasks. Proceeding on the reachable subset would produce a number that looks like the deliverable and does not carry its meaning. Unblocking it requires the authors’ image release, which is outside our control and slow, and which should be requested now if it has not been.

5.6 C5 — Downstream augmentation (optional)

What the plan asked. Run only if C4 is clearly positive.

Not run, and why. Its precondition was not met.

6 STAGE D — FEW-SHOT FACE GENERATION DISPOSITION

What the plan asked. No new runs. The existing negative result stands and is retained, reinterpreted as a task-family mismatch rather than as evidence against the mechanism, with only the benchmark-valid ten-image setting reported and the larger budgets treated as internal diagnostics.

Result. Complied with by construction: no runs were performed. We record the reinterpretation the plan directs, and note that Stage A supplies independent support for it. The standard setup offers a handful of highly heterogeneous target domains rather than many repeated instances of one reusable relation, and A1 shows the method’s advantage concentrated exactly where a reusable relation carries information that a single target image does not. The limitation to state is that few heterogeneous target domains imply weak reusable task-family structure.

7 CROSS-CUTTING REQUIREMENTS

7.1 Baselines

The plan fixes a minimum baseline set for any headline comparison. Present and reported here: target-only conditioning, direct source reuse, relation-only prediction, additive basis with transport, and transport with and without refinement. Absent: the modulation arm with matched transport, for the reasons in A2; low-rank or adapter adaptation, which the plan itself notes does not exist and would be new work to be scheduled only after Stage A; and full or partial fine-tuning, which is implemented and reachable but was not run.

On the fine-tuning reference. The plan cautions against using equal optimisation steps as the only fairness criterion for it. We add that it is the measurement that would attribute the oracle gap of Section 3.1 either to the task or to the basis and backbone, and identify it as the cheapest informative next step.

7.2 Evaluation discipline

Every condition is paired on the same episodes, target-query samples and generation seeds, and this was strengthened mid-stage when refinement noise was brought under a per-episode stream. All differences are reported as absolute paired differences with 95% intervals. No ratio with a vanishing denominator is headlined anywhere in this report; where such a ratio appears in a source file it is accompanied by the absolute difference, which is the quantity we quote. No raw loss curve is presented as evidence about task specificity.

Table 4: Status of the ten deliverables the plan fixes.

#	Deliverable	Status
1	Gaussian-mixture mechanism and calibration	Complete
2	Source-dependence controls	Complete
3	Additive basis against the modulation arm	Deferred — preconditions 0.1 and 0.3 unmet
4	Target-scarcity curve, crossing named	Complete; the answer is that there is no crossing
5	Source-abundance curve	Complete; saturated from the first point
6	Transport-structure ablation	Complete
7	One demographic scarcity experiment	Deferred — C0 qualification is disabling
8	Paired 95% intervals for every headline	Complete
9	Downstream augmentation (optional)	Deferred — precondition C4 unmet
10	Few-shot face stress test (optional)	Deferred — excluded by the plan

Refinement, and the plan’s rule about it. The plan directs that refinement not be promoted into the headline method unless it shows a reliable generation benefit. It does not. The two instruments disagree on its sign, and four measurements of the one-image cell span 0.023 while their half-widths are approximately 0.014. It stays out. Its value is real when the starting coordinate is far away — reusing the source coordinate improves from 1.439 to 0.521 across the target-budget range through refinement alone — and unmeasurable once transport has landed close.

7.3 Excluded work

The dSprites intermediate stage. Excluded by the plan, its role reassigned to the demographic experiment. Section 3.1 gives that decision more weight than it had when taken: the advantage over target-only modelling disappears by twenty target images precisely because a single target image already reveals the task, which is the condition dSprites was designed to remove. With Stage C blocked at C0, no experiment in this stage supplies that condition.

The severity sweep. Excluded unless required to interpret a Stage A result. It was not required. The flag remains wired and verified.

A sweep of the training-time source budget. We considered whether the flat A4 curve reflects what the transport map learned under a fixed training budget rather than a property of the task. Testing that would require retraining at several source budgets. The plan contains no such experiment — every mention of the source budget in it refers to the evaluation script, and its overall posture is that the implementation is frozen. It is recorded here as an open question and was not run.

8 FINAL DELIVERABLES

9 LIMITATIONS

Four limitations travel with every number here.

First, all CIFAR-100 results come from a single checkpoint and a single seed. Every interval is a within-run interval over episodes and contains no seed-to-seed component. This matters most in A5, where four separately trained models are compared and run-to-run variation is the dominant nuisance term; the floor reported there is estimated from four runs, not bounded.

Second, all readings are on the validation split. The test split is untouched and reserved for a single final number, as the plan requires.

Third, absolute generative quality is unmeasured. The kernel Inception distance of 0.4 was not implemented, and no real-against-real control establishes the floor of the sliced Wasserstein distance at these sample sizes. The oracle row functions as a practical floor, but no absolute distance reported here should be read as a distance from zero. What this report establishes is that the coordinate controls generation, not that the generated images are good.

Fourth, the two instruments disagree in at least one place where the disagreement changes a conclusion, and the resolution in Section 3.1 rests on an argument about what each instrument can see rather than on a third measurement.

10 CONCLUSION

The criterion fixed before the stage began is met, and met narrowly: knowing the source task and the relation beats both relation-only transfer and target-only modelling when target data are scarce, and the margin over target-only modelling is gone by twenty target images. The positive control confirms the mechanism is sound where the ground truth is known and locates the difficulty on images precisely: the well is shallow, not absent. Two supporting results are negative in a useful way — source evidence saturates below the smallest budget measured, so the method is cheap on the source side, and the non-linear transport map is not distinguishable from a linear one, so the architecture can be simplified rather than defended.

Three items would change what can be claimed. The unconstrained fine-tuning reference would establish whether the remaining gap to the oracle belongs to the task or to the architecture. The branch integration of 0.1, a repository operation rather than an experiment, unblocks deliverable 3. And the image release behind C0 is the only route to a task family in which a single target example does not already reveal the task — which, on the evidence of A1, is the condition under which this method has the most to offer.